



# Session 1: Background and Definitions


Yehuda Lindell  
Bar-Ilan University

# Secure Computation in Practice



Bar-Ilan University  
Dept. of Computer Science

- ▶ **A request from 1 month ago:**
  - A nonprofit organization in New York, under contract from the US government is doing research on criminal justice
  - The organization asked the US immigration authorities for the list of “Alien Registration Numbers” of aliens arrested in New York City
    - To see which of them are on their list
  - Neither party can hand over their list due to privacy concerns
- ▶ **This is secure set intersection**



# Secure Multiparty Computation



Bar-Ilan University  
Dept. of Computer Science


- ▶ A set of parties with **private** inputs
- ▶ Parties wish to jointly compute a function of their inputs so that certain security properties are preserved
- ▶ Properties must be ensured even if some of the parties **maliciously** attack the protocol
- ▶ Can model any cryptographic task

# Security Requirements



Bar-Ilan University  
Dept. of Computer Science

- ▶ Consider a secure auction (with secret bids):
  - An adversary may wish to learn the bids of all parties – to prevent this, require **PRIVACY**
  - An adversary may wish to win with a lower bid than the highest – to prevent this, require **CORRECTNESS**
  - But, the adversary may also wish to ensure that it always gives the highest bid – to prevent this, require **INDEPENDENCE OF INPUTS**
  - An adversary may try to abort the execution if its bid is not the highest – require **FAIRNESS**




# General Security Properties



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Privacy:** only the output is revealed
- ▶ **Correctness:** the function is computed correctly
- ▶ **Independence of inputs:** parties cannot choose inputs based on others' inputs
- ▶ **Fairness:** if one party receives output, all receive output
- ▶ **Guaranteed output delivery**






# Defining Security



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Option 1: analyze security concerns for each specific problem**
  - Auctions: as in previous slide
  - Elections: privacy, correctness and fairness only (?)
- ▶ **Problems:**
  - How do we know that all concerns are covered?
  - Definitions are application dependent and need to be redefined from scratch for each task



# Defining Security



Bar-Ilan University  
Dept. of Computer Science

- ▶ Option 2: **general definition** that captures all (most) secure computation tasks
  
- ▶ Properties of any such definition
  - Well-defined adversary model
  - Well-defined execution setting
  - Security guarantees are clear and simple to understand

# Modeling Adversaries




Bar-Ilan University  
Dept. of Computer Science

## ▶ Adversarial behavior

- **Semi-honest:** follows the protocol specification
  - Tries to learn more than allowed by inspecting transcript
- **Malicious:** follows any arbitrary strategy
- **Covert:** follows any arbitrary strategy, but is averse to being caught...

## ▶ Adversarial power

- **Polynomial-time**
- **Computationally unbounded:** information-theoretic security






# Modeling Adversaries



Bar-Ilan University  
Dept. of Computer Science

## ► Corruption strategy

- **Static:** the set of corrupted parties is fixed before the execution begins
- **Adaptive:** the adversary can corrupt parties during the execution, based on what has happened
  - Models modern “hacking”
  - Cannot use strategies that choose a small set of representatives to compute for all
  - In general, much harder!



# Execution Setting



Bar-Ilan University  
Dept. of Computer Science

## ► Stand-alone


- Consider a single protocol execution only (or that only a single execution is under attack)

## ► Concurrent general composition

- Arbitrary protocols executed concurrently
- Realistic setting, very important model

## ► Stand-alone vs composition

- Stand-alone: a good place to start studying secure computation, techniques and tools are helpful
- Composition: true goal for constructions




# Feasibility of Secure Computation



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Assuming an honest majority, any functionality can be securely computed**
  - Even information theoretically, and with adaptive security
- ▶ **Without an honest majority, it is impossible to achieve fairness in general**
  - Intuition behind proof of impossibility – later
  - Current understanding of fairness
- ▶ **Without an honest majority, any funct. can be securely computed without fairness**



# Preliminaries




Bar-Ilan University  
Dept. of Computer Science

## ▶ Notations:

- Security parameter  $n$
- We wish security to hold for all inputs of all lengths, as long as  $n$  is large enough

- ## ▶ Function $\mu$ is negligible:
- if for every polynomial  $p(\cdot)$  there exists an  $N$  such that for all  $n > N$  we have  $\mu(n) < 1/p(n)$



# Preliminaries



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Probability ensemble  $X=\{X(a,n)\}$** 
  - Infinite series, indexed by a string  $a$  and natural  $n$
  - Each  $X(a,n)$  is a random variable
    - In our context: output of protocol execution with input  $a$  and security parameter  $n$
    - Probability space: randomness of parties



# Preliminaries



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Computational indistinguishability  $X \approx Y$** 
  - For every (non-uniform) polynomial-time distinguisher  $D$  there exists a negligible function  $\mu$  such that for every  $a$  and all large enough  $n$ 's:
$$|\Pr[D(X(a,n)=1)] - \Pr[D(Y(a,n)=1)]| < \mu(n)$$
- ▶ **Statistical closeness**
  - The same but  $D$  is unbounded in running time

# Notation



Bar-Ilan University  
Dept. of Computer Science

## ► Functionality

- $f=(f_1,\dots,f_m)$ : for input vector  $x$ , each  $f_i(x)$  is a random variable (for probabilistic functionalities)
- Party  $P_i$  receives  $f_i$
- We denote  $(x,y) \rightarrow (f_1(x,y),f_2(x,y))$

# Semi-Honest Adversaries



Bar-Ilan University  
Dept. of Computer Science

## ► Simulation:

- Given input and output, can generate the adversary's view of a protocol execution
- Important: since parties follow protocol, the inputs are well defined

# Semi-Honest Adversaries

- ▶ For every semi-honest  $A$ , there exists a simulator  $S$  such that for every set of corrupted parties  $I$  and every vector of inputs  $x$ , the following are *close*
  - The output of  $A$ , and the outputs of all parties after a protocol execution
  - The output of  $S$  given  $x_i$  and  $f_i(x)$  for all  $i \in I$ , and all the values  $f_1(x), \dots, f_m(x)$


# Security Levels



Bar-Ilan University  
Dept. of Computer Science

## ► Defining “close”

- Computational security = computational indistinguishability
- Statistical security = statistical closeness
- Perfect security = identical distributions






# Semi-Honest Adversaries




Bar-Ilan University  
Dept. of Computer Science



# Semi-Honest Adversaries



Bar-Ilan University  
Dept. of Computer Science




# Properties



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Correctness, independence of inputs, fairness are all non-issues in the semi-honest model**
- ▶ **Why is privacy guaranteed by this definition?**
  - The adversary's view in an execution can be generated from the input and output only
  - If the adversary can compute something after a real protocol execution, it can compute it just from the input/output
  - Very similar to zero-knowledge




# Joint Distribution



Bar-Ilan University  
Dept. of Computer Science

- ▶ A crucial point: need to consider the **joint distribution** of adversary's output and honest parties' output
- ▶ In the definition:
  - We compare the distribution of all inputs and outputs together with the adversary's output




# Joint Distribution



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Example:**
  - **Functionality:** A outputs random bit, B outputs nothing
    - B should clearly not learn A's output bit
  - **Protocol:** A chooses a random bit, outputs it, and sends the bit to B (who ignores it)
  
- ▶ **This is simulatable when separately looking at distribution of B's view and actual outputs**






# Deterministic Functionalities



Bar-Ilan University  
Dept. of Computer Science

- ▶ In the case of deterministic functionalities, the outputs are fully determined by the inputs
- ▶ It suffices to separately prove
  - Correctness
  - Simulation: can generate view of semi-honest adversary (corrupted parties' view), given inputs and outputs only
    - This is significantly easier!




# Malicious Adversaries



Bar-Ilan University  
Dept. of Computer Science

- ▶ First attempt: require the existence of a simulator that generates the adversary's view given the inputs/outputs of corrupted
- ▶ Problem: what are the inputs used by the adversary?
  - They are not necessarily those written on the input tape
  - They are not explicit: the adversary doesn't run the protocol but arbitrary code




# Malicious Adversaries



Bar-Ilan University  
Dept. of Computer Science

- ▶ We also need to require independence of inputs, correctness, fairness etc.
  - These properties are not captured by “view simulation” alone
- ▶ Can we separate correctness and privacy?
  - Instead of computing  $f$ , compute a function that reveals first input bit of other party
  - Correctness or privacy???
- ▶ What about independence of inputs and privacy?




# The Ideal/Real Paradigm



Bar-Ilan University  
Dept. of Computer Science

- ▶ **What is the best we could hope for?**
  - An incorruptible trusted party
  - All parties send inputs to trusted party (over perfectly secure communication lines)
  - Trusted party computes output
  - Trusted party sends each party its output (over perfectly secure communication lines)
  - This is an **ideal world**
- ▶ **What can an adversary do?**
  - Just choose its input...




# The Ideal/Real Paradigm



Bar-Ilan University  
Dept. of Computer Science

- ▶ The real protocol must be like the ideal world
- ▶ Formalizing this notion:
  - For *every adversary  $A$*  attacking the real protocol, *there exists an adversary  $S$*  in the ideal model such that the output distributions (of all) are close
    - Computational indistinguishability, statistical closeness or identical distributions...
  - $S$  simulates a real protocol execution while interacting in the ideal world
  - Here we always look at the joint output distribution

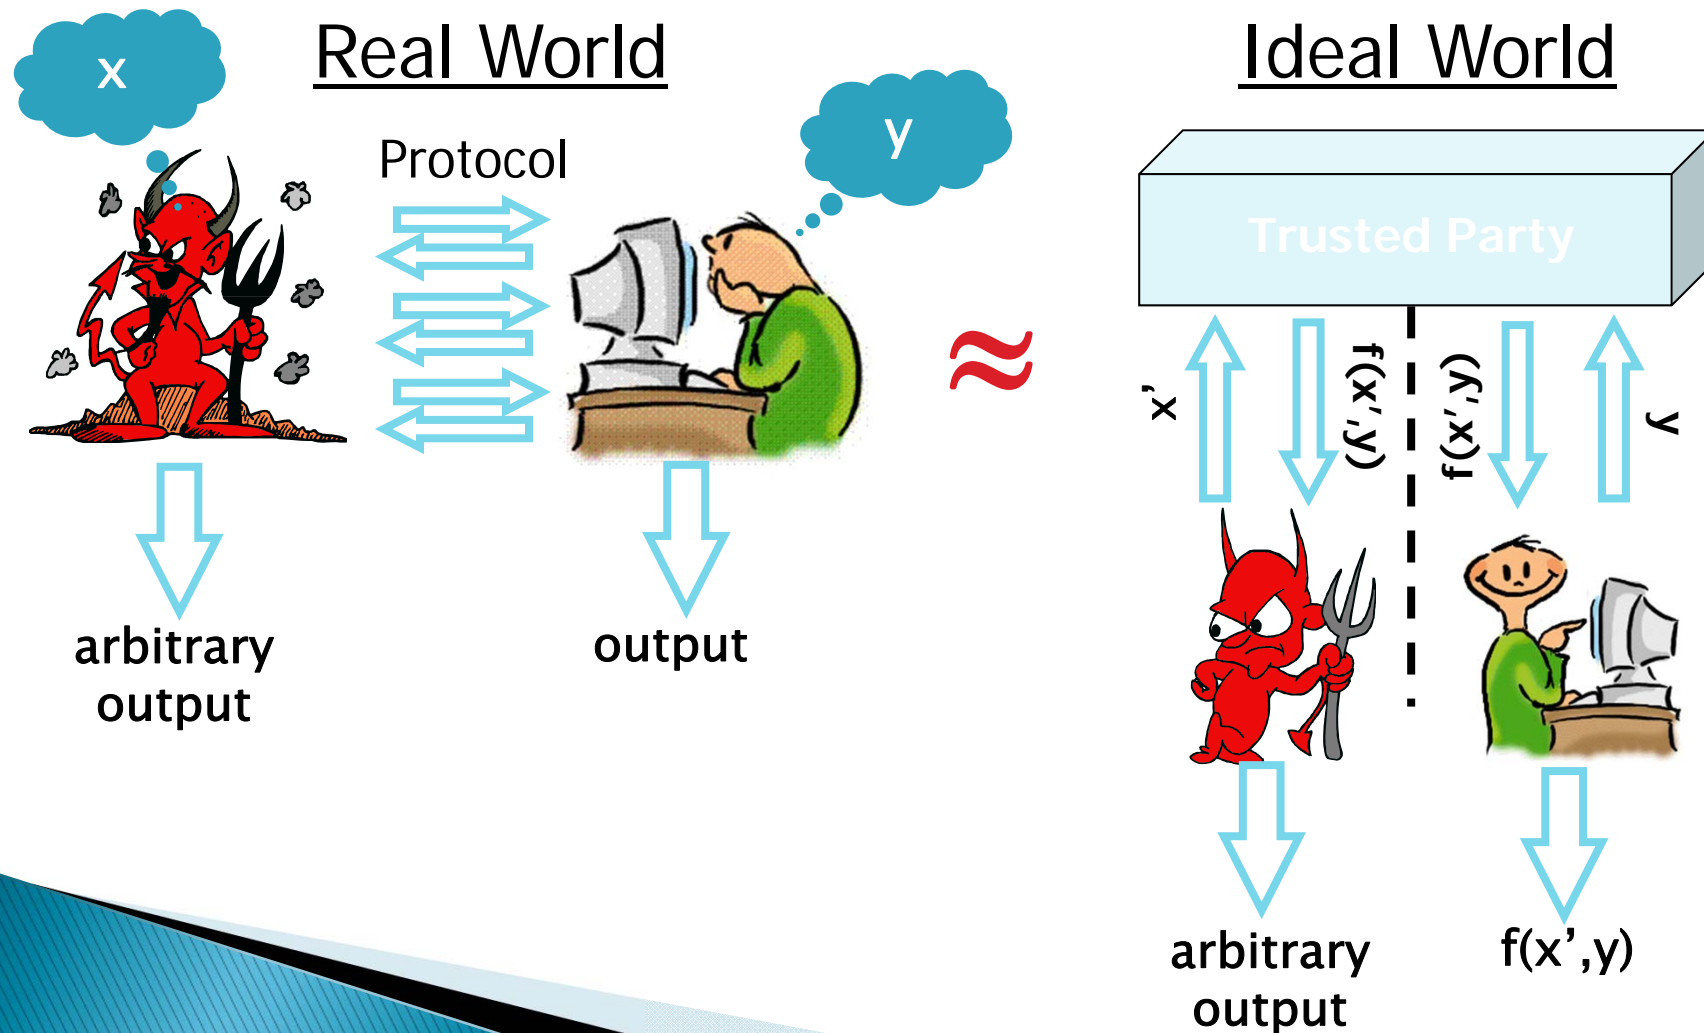




# The Ideal/Real Paradigm



Bar-Ilan University  
Dept. of Computer Science




# “Formal” Security Definition



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Protocol  $\pi$  securely computes a function  $f$  if:**
  - For every non-uniform polynomial-time real-model adversary  $A$ , there exists a non-uniform polynomial-time ideal-model adversary  $S$ , such that for all input vectors and auxiliary inputs:
  - the joint outputs of  $A$  and the honest parties in a real execution of  $\pi$  is indistinguishable<sup>\*</sup> from the joint outputs of  $S$  and the honest parties in an ideal execution where the trusted party computes  $f$

\* Computationally indistinguishable,  
statistically close or identical distributions for  
computational, statistical and perfect security




# Properties



Bar-Ilan University  
Dept. of Computer Science

- ▶ **The following properties hold**
  - Privacy: from adversary's outputs
  - Correctness: from honest parties' outputs
  - Independence of inputs: from ideal execution
  - Fairness and guaranteed output delivery: from ideal execution
  
- More?




# Relaxing the Ideal Model



Bar-Ilan University  
Dept. of Computer Science

- ▶ In some cases, this ideal model is too strong and cannot be achieved
- ▶ Fairness cannot be achieved in general without an honest majority
  - Consider two parties and consider removing the last message of the protocol execution
    - Works for coin tossing...




# Relaxing the Ideal Model



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Change the instructions of the trusted party**
  - Trusted party receives input from all parties
  - Trusted party sends corrupted parties' outputs to adversary
  - Adversary says “continue” or “halt”
  - If “continue”, trusted party sends output to honest parties; else, it sends “abort”




# Reactive Functionalities



Bar-Ilan University  
Dept. of Computer Science

- ▶ Functionalities that obtain inputs and provide outputs in stages
- ▶ Examples:
  - Mental poker
  - Commitment schemes
- ▶ This is also useful for relaxing ideal functionalities (give side information to  $S$ )
- ▶ The definition extends naturally to this as well







# Advantages of This Approach



Bar-Ilan University  
Dept. of Computer Science

- ▶ General – it captures ALL applications
- ▶ The specifics of an application are defined by its functionality, security is defined as above
- ▶ The security guarantees achieved are easily understood (because the ideal model is easily understood)
  - We can be confident that we did not “miss” any security requirements




# Restricted vs General Functionalities



Bar-Ilan University  
Dept. of Computer Science

- ▶ When constructing protocol for general secure computation, it suffices to consider
  - Deterministic functionalities: to compute a probabilistic functionality  $f$ , define  $g((x,r),(y,s))=f(x,y;r\oplus s)$
  - Single-output functionalities: encrypt and MAC the output of the other party
  - Non-reactive functionalities: to compute a reactive functionality, define a series of functions that input/output shared state information (with a MAC)




# Sequential Modular Composition



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Sequential modular composition:**
  - Secure protocols are run sequentially, with arbitrary messages sent in between them
- ▶ **Why consider this?**
  - An important security goal within itself
  - Very helpful (if not crucial) tool for analyzing the security of protocols
- ▶ **Formalization – Hybrid Model**
  - A trusted party helps to compute a sub-functionality
  - REAL messages & IDEAL messages




# Sequential Modular Composition



Bar-Ilan University  
Dept. of Computer Science

- ▶ Subprotocols  $\rho_i$  securely compute functionalities  $f_i$
- ▶ Protocol  $\pi$  securely computes  $g$  in a hybrid model where a trusted party is used to compute every  $f_i$ 
  - This is much easier to analyze since each  $f_i$  is effectively “perfectly secure”
- ▶ **Theorem:** assuming the above, the real protocol  $\pi^\rho$  that uses real calls to each  $\rho_i$  instead of a trusted party for  $f_i$ , securely computes  $g$ .




# Sequential Modular Composition



Bar-Ilan University  
Dept. of Computer Science

## ▶ Proof Sketch

- Assume that a protocol  $\pi$  with a single call to  $f$  securely computes  $g$
- Assume that  $\pi^\rho$  is not secure; an adversary  $A$  breaks the protocol (with  $D$  that distinguishes real from ideal)
- We construct an adversary  $A'$  and distinguisher  $D'$  to attack  $\rho$
- $A'$  receives as auxiliary input the execution prefix of  $\pi$  until  $\rho$  begins, that matches the inputs given in  $\rho$
- After the execution,  $D'$  receives the outputs of all, and uses the auxiliary input to complete the execution of  $\pi$
- $D'$  runs  $D$  and outputs whatever it does



# Sequential Modular Composition



Bar-Ilan University  
Dept. of Computer Science

## ▶ Proof Sketch

- If  $D'$  received the output of an ideal execution of  $f$ , then the output is the same as  $D$  after an ideal execution of  $g$ 
  - This is by the proof of security of  $\pi$  in the hybrid model
- If  $D'$  received the output of a real execution of  $\rho$ , then the output is the same as  $D$  after a real execution of  $\pi^\rho$
- Since  $D$  distinguishes between ideal- $g$  and real- $\pi^\rho$  it follows that  $D'$  distinguishes between ideal- $f$  and real- $\rho$




# Concurrent Composition



Bar-Ilan University  
Dept. of Computer Science

- ▶ **We have considered the stand-alone model**
  - This implies sequential composition
- ▶ **What about concurrent composition?**
  - An Internet-like setting where many (arbitrary, secure and insecure) protocols are run concurrently, with the adversary controlling the scheduling
- ▶ **This models the real-world setting more accurately**
  - We don't know what the result is of running stand-alone protocols concurrently with related inputs




# Concurrent Composition



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Concurrent general composition**
  - Strictly harder than the stand-alone model
  - *Impossible* without some trusted set-up assumption (like a common reference string)
- ▶ **The UC definition (universal composability) guarantees security in this setting**
  - Efficient UC security is a special challenge...
- ▶ **Recommended to study UC next, after studying the stand-alone setting**




# Relaxed Definitions



Bar-Ilan University  
Dept. of Computer Science

- ▶ In order to achieve high efficiency, sometimes can consider weaker definitions
  - Semi-honest (but this is very weak)
  - Covert adversaries: adversary may be malicious but is guaranteed to be caught cheating with good probability
    - Suitable where adversaries can be penalized for being caught cheating (e.g., business loss)
  - Privacy only (malicious)
    - Problematic...




# Defining Privacy Only



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Defining privacy only is very difficult**
  - No correctness and independence of inputs, but as we have seen it is hard to separate these properties
  - Composition is not guaranteed
- ▶ **Example:**
  - Function  $f$  with the property that for every  $x$ , there exists a  $y$  (denoted  $y_x$ ) such that  $f(x, y_x) = x$
  - If  $P_2$  can input  $y_x$  implicitly, then it can learn  $x$




# Private OT



Bar-Ilan University  
Dept. of Computer Science


- ▶ Oblivious transfer
  - Sender: has two strings  $x_0, x_1$
  - Receiver: has a choice bit  $b$
  - Outputs: sender learns nothing about  $b$ , receiver learns only of  $x_0, x_1$
- ▶ For oblivious transfer, we know how to define privacy only, for **two-round protocols**
  - Fortunately we also have such protocols





# Private OT

- ▶ **Why do 2 rounds help?**
  - Receiver sends one message
  - Sender replies with one message
- ▶ **Privacy for a malicious sender**
  - Just need to prove indistinguishability of receiver's first message when  $b=0$  and when  $b=1$
  - This can be extended to many messages
- ▶ **Privacy for a malicious receiver**
  - First message is generated before seeing anything
  - Require that for every first message, there exists a bit  $b'$  such that receiver learns nothing about  $x_{b'}$






# Semi-Honest vs Malicious



Bar-Ilan University  
Dept. of Computer Science

- ▶ Now to confuse you all...
- ▶ It is clear that any protocol that is secure in the presence of malicious adversaries is secure in the presence of semi-honest adversaries
  - A malicious adversary is stronger, and can always behave semi-honestly...
- ▶ But, the simulator in the ideal model is also stronger
  - It can change its input
- ▶ Does this make a difference?




# Semi-Honest vs Malicious



Bar-Ilan University  
Dept. of Computer Science

- ▶ Consider the AND function where only  $P_2$  receives output
- ▶ Consider the following protocol:
  - $P_1$  sends its input directly to  $P_2$
- ▶ Is the protocol secure?
  - Corrupted  $P_1$  learns nothing and gives its input directly, so clearly secure
  - Semi-honest  $P_2$  learns  $P_1$ 's input which doesn't happen if  $P_2$ 's input is 0  $\Rightarrow$  not secure!
  - Malicious  $P_2$ : in the ideal model, simulator can always give input 1 and simulate  $\Rightarrow$  secure!



# Semi-Honest vs Malicious




Bar-Ilan University  
Dept. of Computer Science

## ► Fixing this absurdity

- Allow a semi-honest adversary to also change its input
- Arguably, this is legitimate (to choose input)
- This is called **augmented semi-honest**
  - Note: this stronger notion is also needed for the GMW compilation (this afternoon)

## ► Theorem:

- Security for malicious adversaries implies security for augmented semi-honest adversaries



# Summary



Bar-Ilan University  
Dept. of Computer Science

- ▶ **Semi-honest: simulator given input/output generates the adversary's view**
  - Probabilistic functionalities – must consider joint distribution of view and outputs
  - Deterministic functionalities: easier, suffices to separately consider correctness and view simulation
- ▶ **Malicious: ideal-real simulation**
- ▶ **Sequential composition**
- ▶ **Advanced topics**
  - Concurrent composition
  - Relaxed definition
  - Semi-honest vs malicious

