



Session 3: Secure Computation in the Multi- Party Setting

Benny Pinkas
Bar-Ilan University

Overview



Bar-Ilan University
Dept. of Computer Science

- ▶ Secure computation for more than two parties, computing **Boolean** circuits.
- ▶ **GMW** (Goldreich–Micali–Wigderson)
 - First, for semi-honest adversaries.
 - Then, general compiler from semi-honest to malicious
 - # rounds depends on circuit depth
 - O. Goldreich, Foundations of Cryptography, Vol. II, Chapter 7.
- ▶ **BMR** (Beaver–Micali–Rogaway)
 - $O(1)$ rounds

The setting

- ▶ Parties P_1, \dots, P_n
- ▶ Inputs x_1, \dots, x_n (bits, but can be easily generalized)
- ▶ Outputs y_1, \dots, y_n

- ▶ The functionality is described as a Boolean circuit.
 - Wlog, uses only XOR (+) and AND gates
 - NOT(x) is computed as a $x+1$
 - Wires are **ordered** so that if wire k is a function of wires i and j , then $i < k$ and $j < k$.

The setting

- ▶ The adversary controls a subset of the parties
 - This subset is defined before the protocol begins (is “non-adaptive”)
 - We will not cover the adaptive case

- ▶ Communication
 - Synchronous
 - Private channels between any pair of parties (can be easily implemented using encryption)

Adversarial models



Bar-Ilan University
Dept. of Computer Science

- ▶ Semi-honest
- ▶ Malicious with no abort
 - GMW: A protocol secure any number of malicious parties
- ▶ Malicious with abort
 - GMW: A protocol secure against a minority of malicious parties with abort (will not be discussed here).

Protocol for semi-honest setting



Bar-Ilan University
Dept. of Computer Science

▶ The protocol:

- Each party shares its input bit
- Scan the circuit gate by gate
 - Input values of gate are shared by the parties
 - Run a protocol computing a sharing of the output value of the gate
 - Repeat
- Publish outputs


Protocol for semi-honest setting

► The protocol:

- Each party shares its input bit
- The sharing procedure:
 - P_i has input bit x_i
 - It chooses random bits $r_{i,j}$ for all $i \neq j$.
 - Sends bit $r_{i,j}$ to P_j .
 - Sets its own share to $r_{i,i} = x_i + (\sum_{j \neq i} r_{i,j}) \bmod 2$
 - Therefore $\sum_{j=1 \dots n} r_{i,j} = x_i \bmod 2$.
- Now every P_j has n shares, one for each input x_i of each P_i .

Evaluating the circuit

- ▶ Scan circuit by the order of wires
- ▶ Wire c is a function of wires a, b
- ▶ P_i has shares a_i, b_i . Must get share of c_i .



▶ Addition gate:

- ▶ P_i computes $c_i = a_i + b_i$.
- ▶ Indeed, $c = a + b \pmod{2} = (a_1 + \dots + a_n) + (b_1 + \dots + b_n) = (a_1 + b_1) + \dots + (a_n + b_n) = c_1 + \dots + c_n$

Evaluating multiplication gates

- ▶ $c = a \cdot b = (a_1 + \dots + a_n) \cdot (b_1 + \dots + b_n) =$
 $\sum_{i=1 \dots n} a_i b_i + \sum_{i \neq j} a_i b_j =$
 $\sum_{i=1 \dots n} a_i b_i + \sum_{1 \leq i < j \leq n} (a_i b_j + a_j b_i)$
- ▶ P_i will obtain a share of $a_i b_i + \sum_{i < j \leq n} (a_i b_j + a_j b_i)$
- ▶ Computing $a_i b_i$ by P_i is easy
- ▶ What about $a_i b_j + a_j b_i$?
- ▶ P_i and P_j run the following protocol for every $i < j$.

Evaluating multiplication gates

- ▶ Input: P_i has a_i, b_i , P_j has a_j, b_j .
- ▶ P_i outputs $a_i b_j + a_j b_i + s_{i,j}$. P_j outputs $s_{i,j}$.
- ▶ P_j :
 - Chooses a random $s_{i,j}$
 - Computes the **four** possible outcomes of $a_i b_j + a_j b_i + s_{i,j}$, depending on the four options for P_i 's inputs.
 - Sets these values to be its input to a **1-out-of-4 OT**
- ▶ P_i is the receiver, with input $2a_i + b_i$.



Recovering the output bits

- ▶ The protocol computes shares of the output wires.
- ▶ Each party sends its share of an output wire to the party P_i that should learn that output.
- ▶ P_i can then sum the shares, obtain the value and output it.

Proof of Security

- ▶ **Recall definition of security for semi-honest setting:**
 - Simulation – Given input and output, can generate the adversary's view of a protocol execution.
- ▶ Suppose that adversary controls the set J of all parties but P_i .
- ▶ The simulator is given (x_j, y_j) for all $P_j \in J$.

The simulator

- ▶ **Shares of input wires:** $\forall j \in J$ choose
 - a random share $r_{j,i}$ to be sent from P_j to P_i ,
 - and a random share $r_{i,j}$ to be sent from P_i to P_j .
- ▶ **Shares of multiplication gate wires:**
 - $\forall j < i$, choose a **random** bit as the value learned in the 1-out-of-4 OT.
 - $\forall j > i$, choose a random $s_{i,j}$, and set the four inputs of the OT accordingly.
- ▶ **Output wire y_j of $j \in J$:** set the message received from P_i as the XOR of y_j and the shares of that wire held by $P_j \in J$.

Security proof



Bar-Ilan University
Dept. of Computer Science

- ▶ **The output of the simulation is distributed identically to the view in the real protocol**
 - Certainly true for the random shares $r_{i,j}$, $r_{j,i}$ sent from and to P_i .
 - OT for $j < i$: output is random, as in the real protocol.
 - OT for $j < i$: input to the OT defined as in the real protocol.
 - Output wires: message from P_i distributed as in the real protocol.

▶ **QED**

Performance



Bar-Ilan University
Dept. of Computer Science

- ▶ **Must run an OT for every multiplication gate**
 - Namely, public key operations per multiplication gate
 - Need a communication round between all parties per every multiplication gate
- Can process together a set of multiplication gates if all their input wires are already shared
- Therefore number of rounds is $O(d)$, where d is the depth of the circuit (counting only multiplication gates).

The BMR protocol


- ▶ Beaver–Micali–Rogaway
 - ▶ A multi–party version of Yao’s protocol
 - ▶ Works in $O(1)$ communication rounds, regardless of the depth of the Boolean circuit.
-
- D. Beaver, S. Micali and P. Rogaway, “The round complexity of secure protocols”, 1990.
 - A. Ben–David, N. Nisan and B. Pinkas, “FairplayMP – A System for Secure Multi–Party Computation”, 2010.

The BMR protocol

- ▶ Two random seeds (garbled values) are set for every wire of the Boolean circuit:
 - Each seed is a concatenation of seeds generated by all players and secretly shared among them.
- ▶ The parties **securely compute together** a 4×1 table for every gate (in parallel):
 - Given 0/1 seeds of the input wires, the table reveals the seed of the resulting value of the output wire.

The BMR protocol

- ▶ The parties **securely compute together** a 4×1 table for every gate (in parallel):
 - This is essentially a secure computation of the table
 - But all tables can be computed in parallel. Therefore $O(1)$ rounds.
 - This is the main bottleneck of the BMR protocol.
- ▶ Given the tables, and seeds of the input values, it is easy to compute the circuit output.



The malicious case

- ▶ What can go wrong with malicious behavior?
 - Using shares other than those defined by the protocol, using arbitrary inputs to the OT protocol and sending wrong shares of output wires...
- ▶ We will show a compiler which forces the parties to operate as in the semi-honest model. (For both GMW and BMR.)
- ▶ The basic idea:
 - In every step, each P_i proves in zero knowledge that its messages were computed according to the protocol

Zero knowledge

(more on this tomorrow)



Bar-Ilan University
Dept. of Computer Science

- ▶ Prover P , verifier V , language L
- ▶ P proves that $x \in L$ without revealing anything
 - **Completeness:** V always accepts when $x \in L$, and an honest P and V interact.
 - **Soundness:** V accepts with negligible probability when $x \notin L$, for any P^* .
 - Computational soundness: only holds when P^* is polynomial-time
- ▶ **Zero-knowledge:**
 - There exists a **simulator** S such that $S(x)$ is indistinguishable from a real proof execution.

A warm-up

- ▶ Assume that each P_i runs a **deterministic** program Π_i . The compiler is the following:
 - Each P_i **commits** to its input x_i by sending $C_i(r_i, x_i)$, where r_i is a random string used for the commitment.
 - Let T_i^s be the **transcript** of P_i at step s , i.e. all messages received and sent by P_i until that step.
 - Define the language $L_i = \{T_i^s \text{ s.t. } \exists x_i, r_i \text{ so that all messages sent by } P_i \text{ until step } s \text{ are the output of } \Pi_i \text{ applied to } x_i, r_i \text{ and to all messages received by } P_i \text{ up to that step}\}$
 - When sending a message in step s prove in zero-knowledge that $T_i^s \in L_i$.

Handling randomized protocols

- ▶ The previous construction assumes that P_i 's program, Π_i , is deterministic.
- ▶ This is not true in the semi-honest protocol we have seen.
 - In particular, the choice of shares, and the sender's input to the OT, must be random.
 - The compiler must ensure that P_i chooses its random coins independently of the messages received from other parties.
 - This is not ensured by the previous construction.

The compiler

- ▶ We will describe **the basic issues** of a protocol secure against any number of malicious parties, but with no aborts allowed.
- ▶ Communication model:
 - Messages are published on a **bulletin board**, and can be read by all parties.
 - This implements a **broadcast**, ensuring that all parties receive the same message.
 - Broadcast can be easily implemented if a **public key infrastructure** exists.
 - We assume that a PKI does exist.

The compiler



Bar-Ilan University
Dept. of Computer Science

- ▶ **Input commitment phase:**
 - Each party commits to its input.
- ▶ **Coin generation phase:**
 - The parties generate random tapes for each other.
 - Initial idea: random tape of P_i is defined as $s_{1,i} \oplus s_{2,i} \oplus \dots \oplus s_{n,i}$, where $s_{j,i}$ is chosen by P_j .
 - But this lets P_n control the outcome ☹
- ▶ **Protocol emulation phase:**
 - Run the protocol while proving that parties operations comply with their inputs and random tapes.

The protocol:

Input commitment phase



Bar-Ilan University
Dept. of Computer Science

- ▶ The required functionality for P_1 is $(x, 1^{|x|}, \dots, 1^{|x|}) \rightarrow (r, C_r(x), \dots, C_r(x))$, and similarly for each P_i .
- ▶ It is not sufficient to ask P_i to just broadcast a commitment of its input
 - This does not ensure that this is a random commitment for which P_i knows a decommitment.
- ▶ The protocol is more complex...
- ▶ It is useful to first design tools that can help in constructing the compiler.

Tool 1: image transmission

- ▶ The required functionality is
 $(a, 1^{|a|}, \dots, 1^{|a|}) \rightarrow (\lambda, f(a), \dots, f(a))$ (all receive the same function of a)
- ▶ Protocol
 - P_1 **broadcasts** an encryption of $f(a)$
 - For $j=2\dots n$, P_1 proves to P_j a zero-knowledge strong proof of knowledge of a value a corresponding to $f(a)$.
 - If P_j **rejects**, it **broadcasts** the coins it used in the proof.
- ▶ **Output:** For $j=2\dots n$, if P_j sees a **justifiable rejection** it **aborts**, otherwise it outputs $f(a)$.

Tool 1: image transmission

- ▶ The required functionality is
$$(a, 1^{|a|}, \dots, 1^{|a|}) \rightarrow (\lambda, f(a), \dots, f(a))$$
- ▶ Agreement as to whether P_1 misbehaved is reduced to the decision on whether some verifier has justifiably rejected the proof.
- ▶ If P_1 is honest, then no malicious party can claim that it cheated.

Tool 2: authenticated computation

- ▶ The required functionality is
 $(a, b_2, \dots, b_n) \rightarrow (\lambda, v_2, \dots, v_n)$, where $v_j = f(a)$ if $b_j = h(a)$ and $v_j = \lambda$ otherwise.
- ▶ Protocol:
 - Use the **image transmission** tool to **broadcast** $(f(a), h(a))$ to all P_j , $j=2 \dots n$.
 - P_j outputs $f(a)$ if $v_j = h(a)$, and λ otherwise.
- ▶ Comment: P_j learns a function $f(a)$ of a , if it already has the function $h(a)$ (e.g., if it has a commitment to a)

Tool 3: multi-party augmented coin-tossing



Bar-Ilan University
Dept. of Computer Science

- ▶ The required functionality is $(1^n, \dots, 1^n) \rightarrow (r, g(r), \dots, g(r))$.
- ▶ Typically we will use it for computing $(1^n, \dots, 1^n) \rightarrow ((r, s), C_s(r), \dots, C_s(r))$.
- ▶ **The challenge:** ensuring that P_1 's output is random. We cannot trust P_1 to choose a random output.

Tool 3: multi-party augmented coin-tossing



Bar-Ilan University
Dept. of Computer Science

- ▶ $(1^n, \dots, 1^n) \rightarrow ((r, s), C_s(r), \dots, C_s(r))$.
 - Toss and commit: $\forall i$, P_i chooses r_i, s_i and uses the image transmission tool to send $c_i = C_{s_i}(r_i)$ to all P_j .
 - Open commits: $\forall i \geq 2$, P_i uses the authenticated computation tool to send s_i, r_i to all parties that already have c_i .
 - If P_j obtains r_i agreeing with c_i , it sets $r_i^j = r_i$ (also, $r_j^j = r_j$). Otherwise it aborts.
 - If P_1 did not abort, it sets $r = \bigoplus_{i=1 \dots n} r_i$ sends $C_s(r)$ to all other parties, and proves that it was constructed correctly.

Tool 3: multi-party augmented coin-tossing (contd.)



Bar-Ilan University
Dept. of Computer Science

- ▶ P_1 sends $C_s(r)$ to all other parties, and proves that it was constructed correctly.
- ▶ Run the **authenticated computation** functionality
 - ▶ P_1 chooses a random s . Its input to the protocol is $(r_1, s_1, s, \oplus_{j=2 \dots n} r_i^1)$
 - ▶ P_j 's input is $c_1, \oplus_{j=2 \dots n} r_i^j$.
 - ▶ If $c_1 = C_{s1}(r_1)$ and $\oplus_{j=2 \dots n} r_i^j = \oplus_{j=2 \dots n} r_i^1$, then P_j outputs $C_s(\oplus_{j=1 \dots n} r_i) = C_s(r)$. Otherwise it aborts.
 - ▶ P_1 outputs r .

The main protocol:

Input commitment phase



Bar-Ilan University
Dept. of Computer Science

► Protocol:

- P_i chooses random r'_i and uses image transmission functionality to send $c' = C_{r'_i}(x_i)$ to all parties.
- Run augmented coin-tossing protocol s.t. P_i learns (r_i, r''_i) and others learn $c'' = C_{r''_i}(r_i)$.
- Run authenticated computation where P_i has input (x_i, r_i, r'_i, r''_i) and others input (c', c'') , and others learn $C_{r_i}(x_i)$ if (c', c'') are the required functions of P_i 's input.

The main protocol: coin generation phase



Bar-Ilan University
Dept. of Computer Science

- ▶ Each P_i runs the augmented coin tossing protocol where
 - P_i learns (r^i, s^i)
 - The other parties learn $C_{s^i}(r^i)$.

The main protocol:

Protocol emulation phase



Bar-Ilan University
Dept. of Computer Science

- ▶ The parties use the authenticated computation functionality
 - $(a, b_2, \dots, b_n) \rightarrow (\lambda, v_2, \dots, v_n)$, where $v_j = f(a)$ if $b_j = h(a)$ and $v_j = \lambda$ otherwise.
- ▶ Suppose that it is P_i 's turn to send a message
 - Its input is (x_i, r_i, T_t) , as well as the coins used for commitments, where T_t is the sequence of messages exchanged so far.
 - Every other party has input $(C(x_i), C(r_i), T_t)$
 - $f(x_i, r_i, T_t)$ is the message P_i must send
 - It is accepted if $(C(x_i), C(r_i), T)$ agree with x_i, r_i, T and the program that is run

Summary



Bar-Ilan University
Dept. of Computer Science

- ▶ Can compute any functionality securely in presence of semi-honest adversaries
- ▶ Protocol is efficient enough for use, for circuits that are not too large
- ▶ Recommendation: read full proof (Goldreich's book).

