




LARGE-SCALE SECURE COMPUTATION

MPC FOR PARALLEL RAM PROGRAMS

Elette Boyle
Technion


Based on joint works with Kai-Min Chung and Rafael Pass

Multi-Party Computation (MPC)



[GMW87] – Computational Setting
[BGW88, CCD88] – Information Theoretic Setting with Secure Channels


This Talk: Large-Scale MPC



[GMW87] – Computational Setting
[BGW88, CCD88] – Information Theoretic Setting with Secure Channels

This Talk: Large-Scale MPC

Lots of Parties



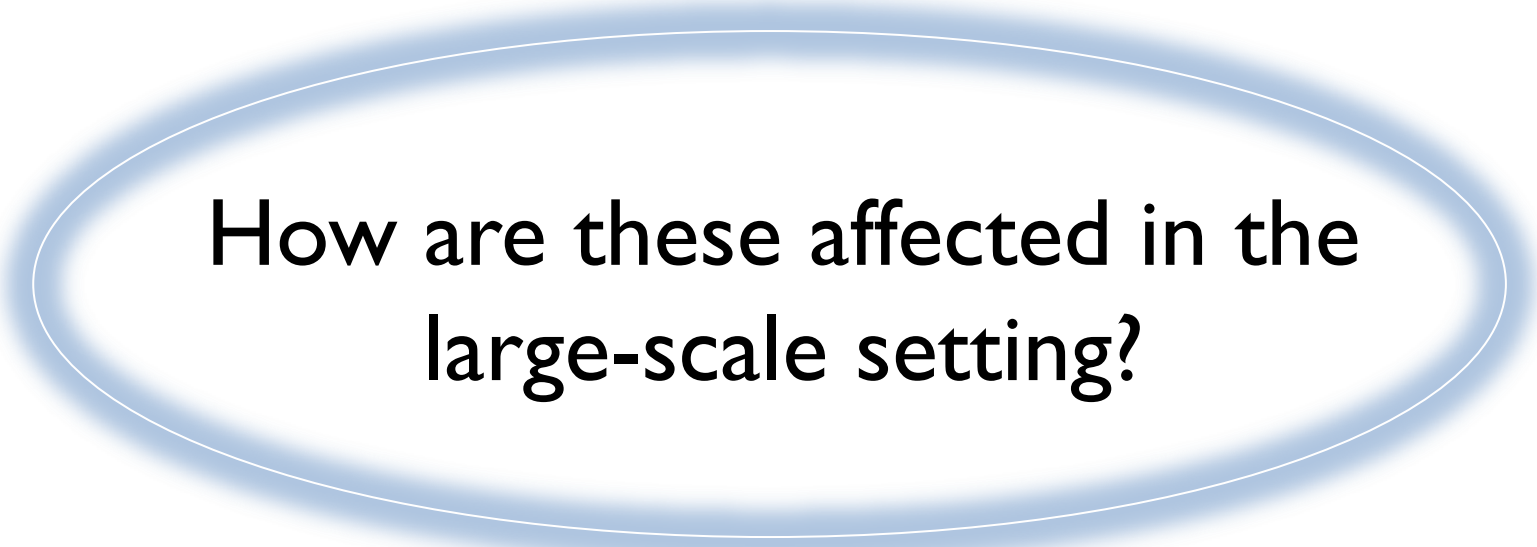
Lots of Data

MPC Efficiency Metrics

Communication

Memory

Computation



How are these affected in the
large-scale setting?

Costs of Communication

- # of bits communicated
- # of sequential rounds
- ...Who a party is speaking to


Nearly all protocols:

Every party speaks to every party

Communication: *Locality* Metric

[BGT13]


parties:
 $n = 5$



Communication: *Locality* Metric

[BGT13]

parties:
 $n = 10,000$



x_i


⋮

Not practical!

Communication: *Locality* Metric

[BGT13]

parties:
 $n = 10,000$



Communication

Locality: Total # parties

each party communicates with throughout protocol lifetime



Memory: Balancing the Burden



- Combined data size is huge!
- Want: Memory requirement per party
$$\approx (\text{his input} + \text{Space}(\Pi)/n)$$



Computation: Going Beyond Circuits

 f 

MPC protocol
for securely
computing the
circuit


Optimizing this transformation
yields better MPC efficiency

Program

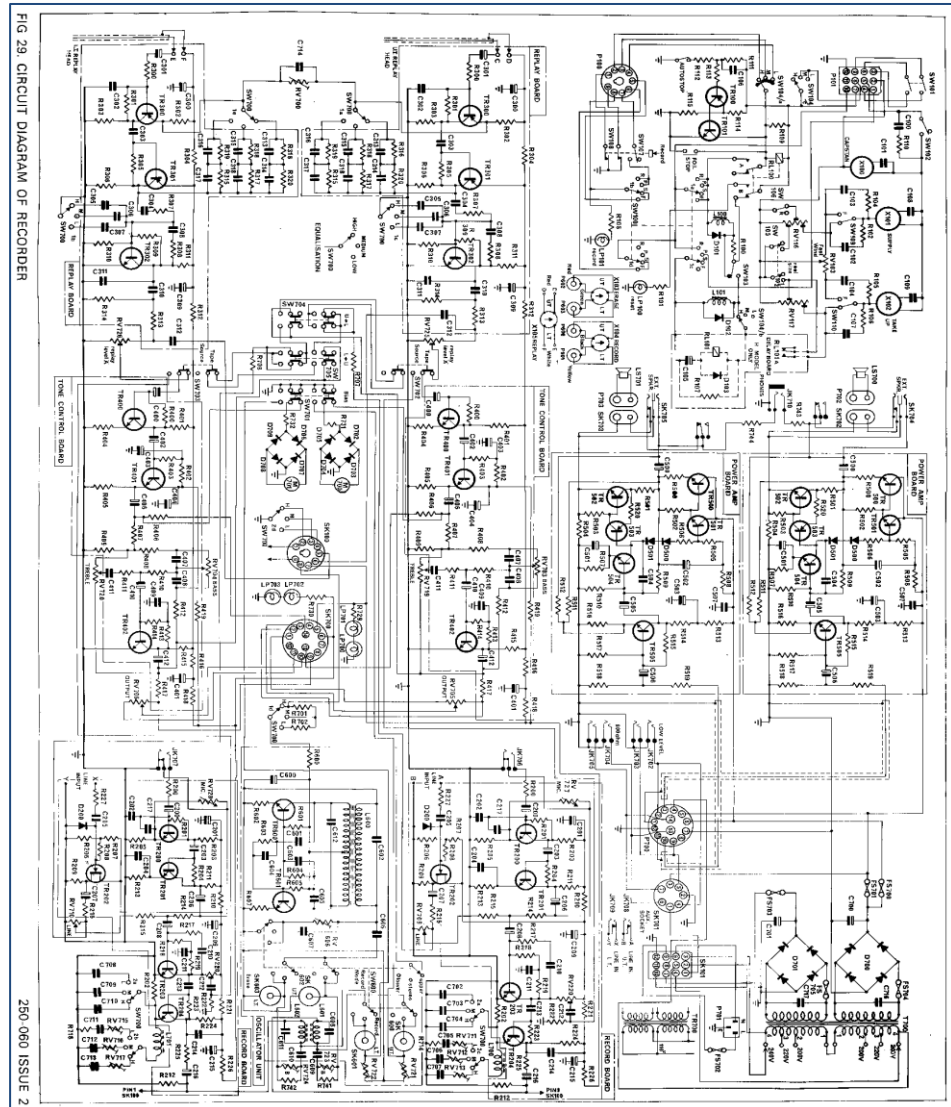
```

1 low = p
2 high = max(p, r + 1)
3 while low < high
4     mid = (low + high) / 2
5     if x ≤ T[mid]
6         high = mid
7     else low = mid + 1
8 return high

```




Generically:
Blow up by factor of
entire database size!





Models of Computation I

- Circuits




AND, OR, NOT gates


- Turing Machines



- RAM Machines



- Parallel RAM Machines



Computation: Going Beyond Circuits


Large-scale computations *f* leverage
random access and ***parallelism***

- Circuit (and TM) model for *f* not appropriate!

Computation: Going Beyond Circuits

Large-scale computations f leverage
random access and ***parallelism***

- Circuit (and TM) model for f not appropriate!
- RAM model for f loses parallelism!
- **Parallel RAM (PRAM) Model**



Rough History of Prior MPC Work

- Circuits model

E.g.: Original protocols [GMW87, BGW88, CCD88,...], Scalable MPC [DI06, DN07, DIK+08, DIK10, DKMS12, ZMS14], MPC on incomplete networks [CGO10, CGO12], MPC based on FHE / Obfuscation [Gen09, AJL+12, MSS13, GGHR14], Optimized MPC for practice [BNP08, KS08, LPS08, NO09, LPI1, BDOZ11, DPSZ12, NNOS12, LI3, FJN+13, ALSZ13, DZ13, LR14, ZRE15,...]

- RAM model

- 2-PC [OS97, GKK+11, LO13, GGHJ+13, GHRW14, WHHSS14]
- Extensions to MPC [DMN11] **don't scale with n**

- PRAM model (nothing)

Eg: Per-party memory requirement \sim size of *all parties'* inputs

Asymptotically

The Goal: Efficient MPC for PRAM

n -party MPC for PRAMs Π

Time Steps - *Parallel Time*(Π)

Needed for
security

Per-party Computation - *Comp*(Π)/ n + His input

Per-party Memory - His input + *Space*(Π)/ n

Comm Locality - 1

Theorem [BCPI4,BCPI5]:

n -party MPC for PRAMs Π

$$\tilde{O} = \text{polylog}(n)$$

Rounds - $\tilde{O}(\text{Parallel Time}(\Pi))$

Per-party Computation - $\tilde{O}(\text{Comp}(\Pi)/n)$

Per-party Memory - $\tilde{O}(\text{His input} + \text{Space}(\Pi)/n)$

Comm Locality - $\tilde{O}(\text{His input}) + \text{BC} / \text{party}$

Given a 1-time (reusable) preprocessing stage


Static corruptions, $2/3 + \epsilon$ honest parties, Unconditional security

The Construction

For Large Data, Many Parties...

- Step 1: Secret Share inputs across parties

Eg: evaluations of random polynomial st $p(0)=s$ [Sha79]



Problem 1: Everyone talks to everyone

Problem 2: Everyone stores all inputs

- Step 2: Evaluate gate-by-gate on shares
(sometimes with communication)


Problem 3: Computation ~ Circuit Size

Consider a Simpler Problem:

Large Data, Few Parties

- Step 1: Secret Share inputs across parties

Eg: evaluations of random polynomial st $p(0)=s$ [Sha79]



Problem 1: Everyone talks to everyone

Problem 2: Everyone stores all inputs

- Step 2: Evaluate gate-by-gate on shares
(sometimes with communication)


Problem 3: Computation ~ Circuit Size

Consider a Simpler Problem:

Large Data, Few Parties

- Step 1: Secret Share inputs across parties

Eg: evaluations of random polynomial st $p(0)=s$ [Sha79]



Problem 1: Everyone talks to everyone

Problem 2: Everyone stores all inputs

These are ok!

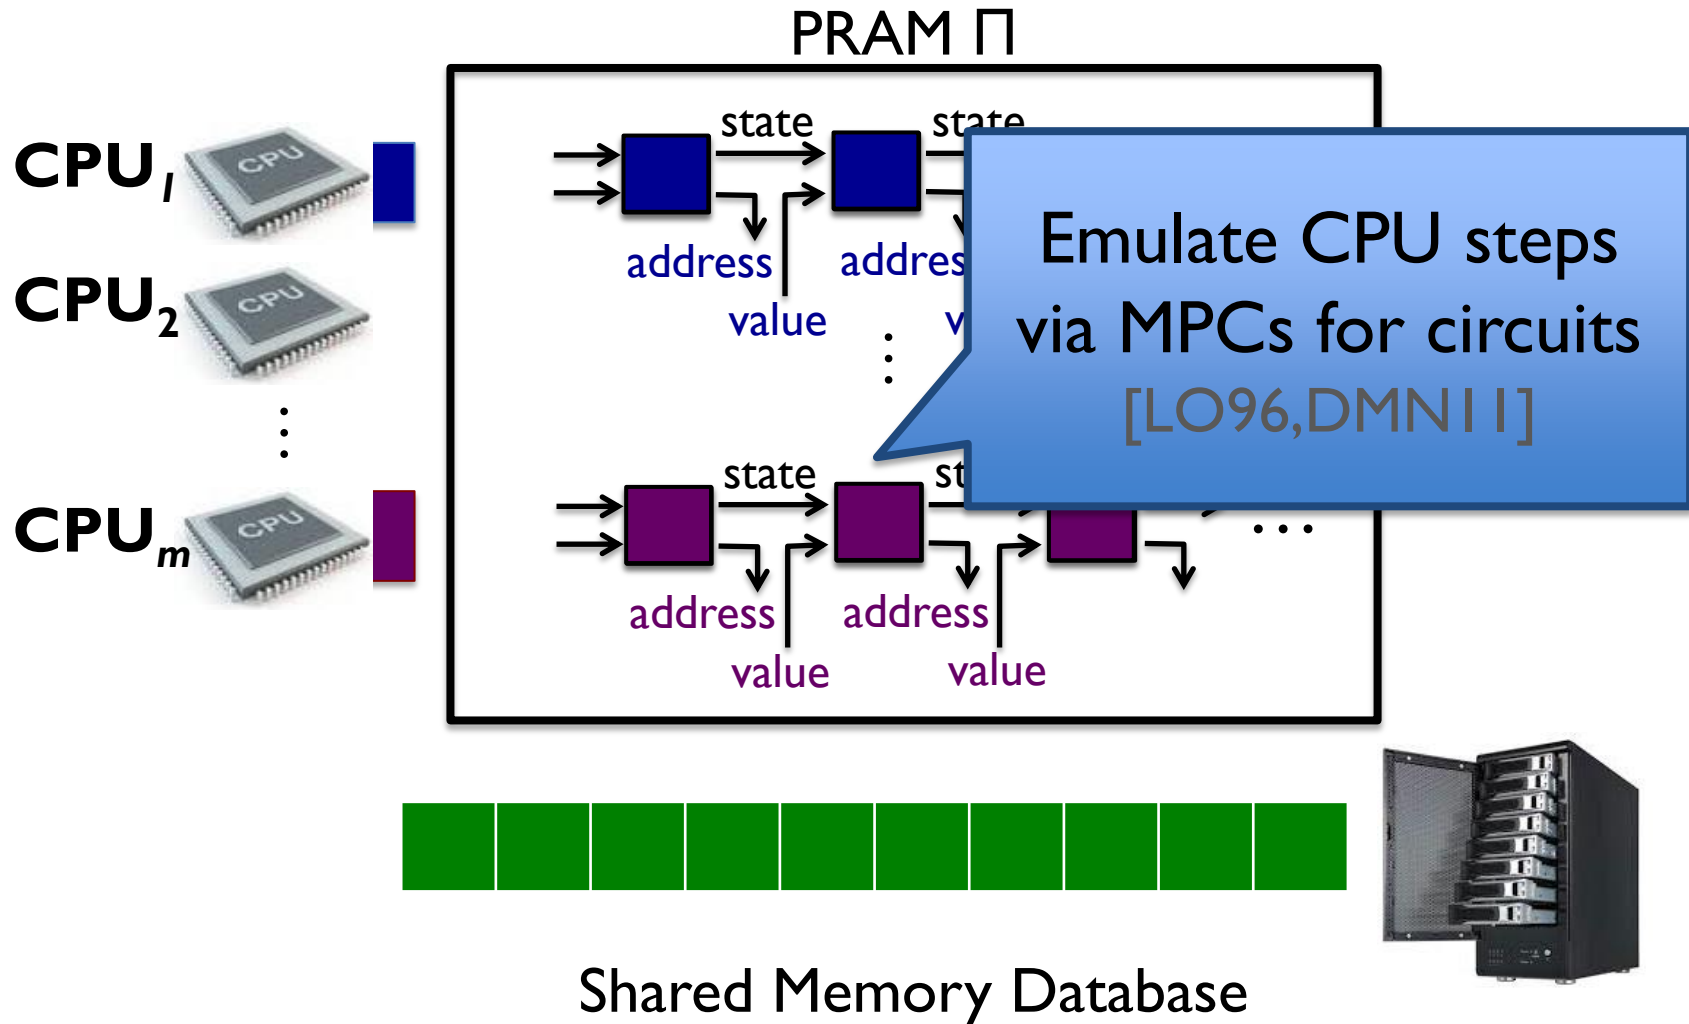
- Step 2: Evaluate gate-by-gate on shares
(sometimes with communication)

Problem 3: Computation ~ Circuit Size

Wanted:
Comp ~ |PRAM|



MPC for PRAM: First Idea



MPC for PRAM: First Idea


à la [LO06, DMNI I]

- Step 1: Secret Share inputs across parties
- Step 2: Emulate PRAM CPU steps via small-scale MPCs



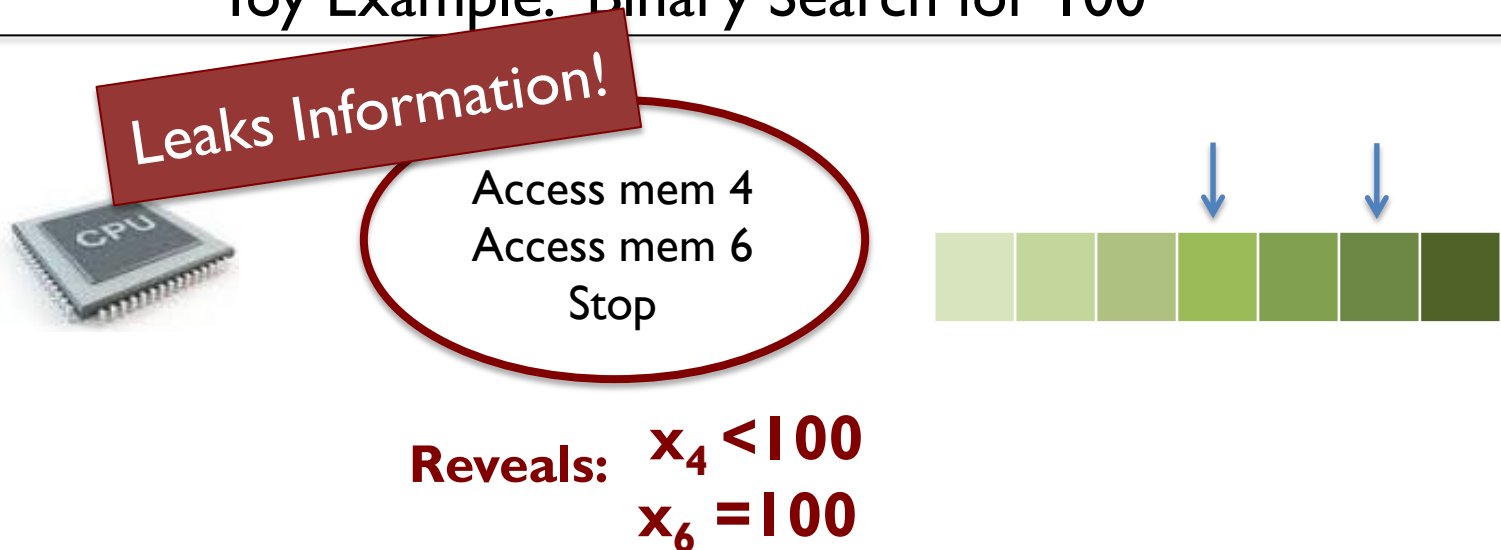
Parties only see addresses & *shares* of secrets!

Addresses may leak information!



Memory Access Patterns May Leak Information!

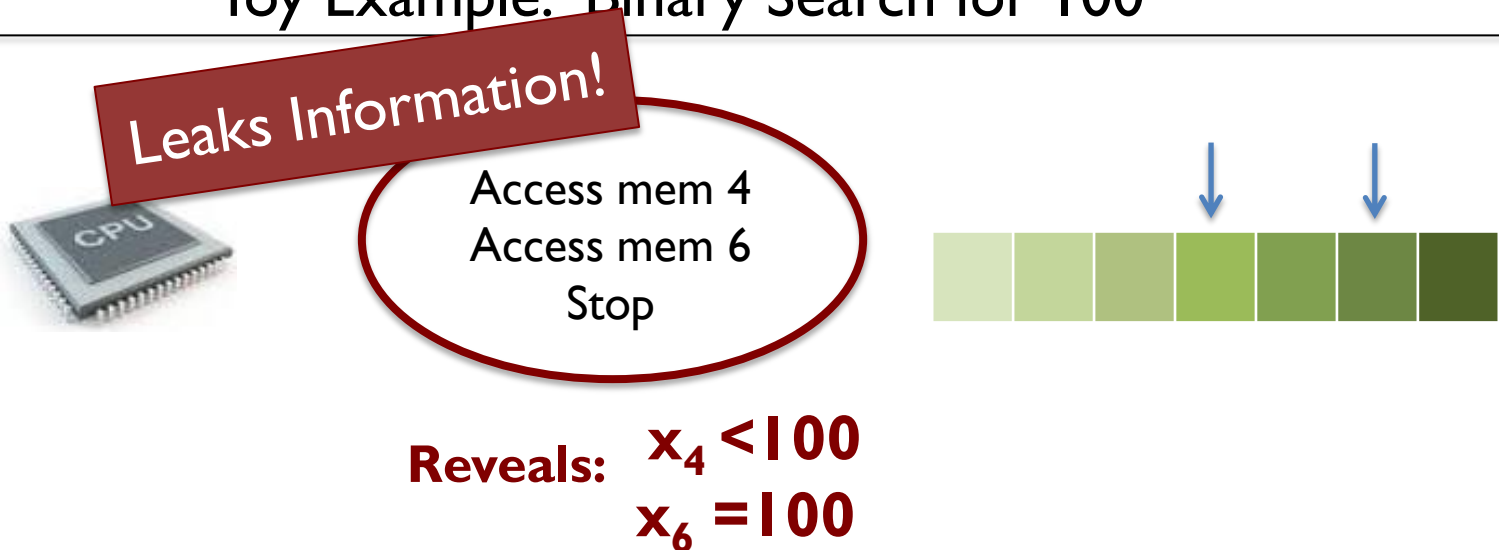
Toy Example: Binary Search for 100




Wanted: PRAM → **Oblivious PRAM**

“Oblivious” = memory access patterns appear independent of data

Toy Example: Binary Search for 100



Oblivious Program Compilers




History:

M = memory size

- Turing Machines: $\log(M)$ overhead [PF 79]
- RAM programs: $\text{polylog}(M)$ overhead [Gol86, Ost90, GO96, Ajt10, DMN11, SCSL11, CPI3, GGHJ+13, SDSF+13]
- PRAM: $\text{polylog}(M)$ overhead [BCPI4]

Core Problem: Supporting Parallel Accesses!



Can't afford for
CPUs to take turns!

Storing multiple copies
causes consistency issues!


Reveals lookup collision!

New Protocol: (Few-Party) **MPC for PRAM**


- Step 1: *Secret Share* inputs across parties



- Step 2: PRAM \rightarrow Oblivious PRAM



- Step 3: Emulate OPRAM via small-scale MPCs



And for **Large Data** and **Many Parties**...


- Step 1: Secret Share inputs across parties

Problem 1: Everyone talks to everyone

Problem 2: Everyone stores all inputs


For another
time...

- Step 2: PRAM \rightarrow Oblivious PRAM ... while load balancing!




- Step 3: Emulate OPRAM via small-scale MPCs

Computation \sim |PRAM|




Teaser of Additional Techniques


Electing Committees




Distributed OPRAM



Load-Balancing via Job Passing



Load-Balanced Routing over Expander Graphs



Future Directions

- “OPRAM is the new ORAM”
- me
- Pushing Large-Scale MPC toward Practicality
 - Leveraging computational assumptions? Adaptive security?
 - Improving broadcast with locality? Honest minority? Targeted protocols?
 - MPC for MapReduce? Asynchronous models?

