

Session 1: Definitions and Oblivious Transfer

Yehuda Lindell
Bar-Ilan University



Secure Multiparty Computation

- A set of parties with **private inputs**
- Parties wish to jointly compute a function of their inputs so that certain security properties are preserved
- Properties must be ensured even if some of the parties **maliciously** attack the protocol
- Can model any cryptographic task

Applications

- **Elections**
- **Auctions**
- **Private database search**
- **Privacy-preserving data mining**
- **Secure set intersection**
- **Much much more...**

Security Requirements

- **Consider a secure auction (with secret bids):**
 - An adversary may wish to learn the bids of all parties
 - to prevent this, require **PRIVACY**
 - An adversary may wish to win with a lower bid than the highest – to prevent this, require **CORRECTNESS**
 - But, the adversary may also wish to ensure that it always gives the highest bid – to prevent this, require **INDEPENDENCE OF INPUTS**
 - An adversary may try to abort the execution if its bid is not the highest – require **FAIRNESS**

General Security Properties

- **Privacy:** only the output is revealed
- **Correctness:** the function is computed correctly
- **Independence of inputs:** parties cannot choose inputs based on others' inputs
- **Fairness:** if one party receives output, all receive output
- **Guaranteed output delivery**

Defining Security

- **Option 1: analyze security concerns for each specific problem**
 - Auctions: as in previous slide
 - Elections: privacy, correctness and fairness only (?)
- **Problems:**
 - How do we know that all concerns are covered?
 - Definitions are application dependent and need to be redefined from scratch for each task

Defining Security

- **Option 2: general definition that captures all (most) secure computation tasks**
- **Properties of any such definition**
 - Well-defined adversary model
 - Well-defined execution setting
 - Security guarantees are clear and simple to understand

Modeling Adversaries

- **Adversarial behavior**
 - **Semi-honest:** follows the protocol specification
 - Tries to learn more than allowed by inspecting transcript
 - **Malicious:** follows any arbitrary strategy
 - **Covert:** follows any arbitrary strategy, but is averse to being caught...
- **Adversarial power**
 - **Polynomial-time:** computational security
 - **Computationally unbounded:** information-theoretic security

Modeling Adversaries

- **Corruption strategy**
 - **Static:** the set of corrupted parties is fixed before the execution begins
 - **Adaptive:** the adversary can corrupt parties during the execution, based on what has happened
 - Models modern “hacking”
 - Cannot use strategies that choose a small set of representatives to compute for all
 - In general, **much harder!**

Execution Setting

- **Stand-alone**
 - Consider a single protocol execution only (or that only a single execution is under attack)
- **Concurrent general composition**
 - Arbitrary protocols executed concurrently
 - Realistic setting, very important model
- **Stand-alone vs composition**
 - Stand-alone: a good place to start studying secure computation, techniques and tools are helpful
 - Composition: true goal for constructions

Feasibility of Secure Computation

- **Assuming an honest majority, any functionality can be securely computed**
 - Even information theoretically, and with adaptive security
- **Without an honest majority, it is impossible to achieve fairness in general**
- **Without an honest majority, any functionality can be securely computed without fairness**

Preliminaries

- **Notations:**
 - Security parameter n
 - We wish security to hold for **all inputs** of all lengths, as long as n is large enough
- **Function μ is negligible:** if for every polynomial $p(n)$ there exists an N such that for all $n > N$ we have $\mu(n) < \frac{1}{p(n)}$

Preliminaries

- **Probability ensemble $X = \{X(a, n)\}$**
 - Infinite series, indexed by a string a and natural n
 - Each $X(a, n)$ is a random variable
 - In our context: output of protocol execution with input a and security parameter n
 - Probability space: randomness of parties

Preliminaries

- **Computational indistinguishability $X \approx Y$**
 - For every (non-uniform) polynomial-time distinguisher D there exists a negligible function μ such that for every a and all large enough n 's:
$$|\Pr[D(X(a, n) = 1)] - \Pr[D(Y(a, n) = 1)]| < \mu(n)$$

Notation

- **Functionality**

- $f = (f_1, \dots, f_m)$: for input vector x , each $f_i(x)$ is a random variable (for probabilistic functionalities)
- Party P_i receives f_i
- We denote $(x, y) \rightarrow (f_1(x, y), f_2(x, y))$

Semi-Honest Adversaries


- **Simulation:**
 - Given input and output, can generate the adversary's view of a protocol execution
 - Important: since parties follow protocol, the inputs are **well defined**

Semi-Honest Adversaries


- For every **semi-honest** A , there exists a **simulator** S such that for every set of corrupted parties I and every vector of inputs x , the following are *computationally indistinguishable*
 - The output of A , and the outputs of all parties after a protocol execution
 - The output of S given x_i and $f_i(x)$ for all $i \in I$, and all the values $f_1(x), \dots, f_m(x)$

Semi-Honest Adversaries

The REAL execution



Simulation



Properties

- **Correctness, independence of inputs, fairness are all non-issues in the semi-honest model**
- **Why is privacy guaranteed by this definition?**
 - The adversary's view in an execution can be generated from the input and output only
 - If the adversary can compute something after a real protocol execution, it can compute it just from the input/output
 - Very similar to zero-knowledge

Joint Distribution

- A crucial point: need to consider the **joint distribution** of adversary's output and honest parties' output
- In the definition:
 - We compare the distribution of all inputs and outputs together with the adversary's output

Joint Distribution

- **Example:**
 - **Functionality:** A outputs random bit, B outputs nothing
 - B should clearly not learn A 's output bit
 - **Protocol:** A chooses a random bit, outputs it, and sends the bit to B (who ignores it)
- **This is simulatable when separately looking at distribution of B 's view and actual outputs**

Deterministic Functionalities

- In the case of deterministic functionalities, the outputs are fully determined by the inputs
- It suffices to separately prove
 - Correctness
 - Simulation: can generate view of semi-honest adversary (corrupted parties' view), given inputs and outputs only
 - This is significantly easier!

Malicious Adversaries

- **First attempt: require the existence of a simulator that generates the adversary's view given the inputs/outputs of corrupted**
- **Problem: what are the inputs used by the adversary?**
 - They are not necessarily those written on the input tape
 - They are not explicit: the adversary doesn't run the protocol but arbitrary code

Malicious Adversaries

- **We also need to require independence of inputs, correctness, fairness etc.**
 - These properties are not captured by “view simulation” alone
- **Can we separate correctness and privacy?**
 - Instead of computing f , compute a function that reveals first input bit of other party
 - Correctness or privacy???
- **What about independence of inputs and privacy?**

The Ideal/Real Paradigm

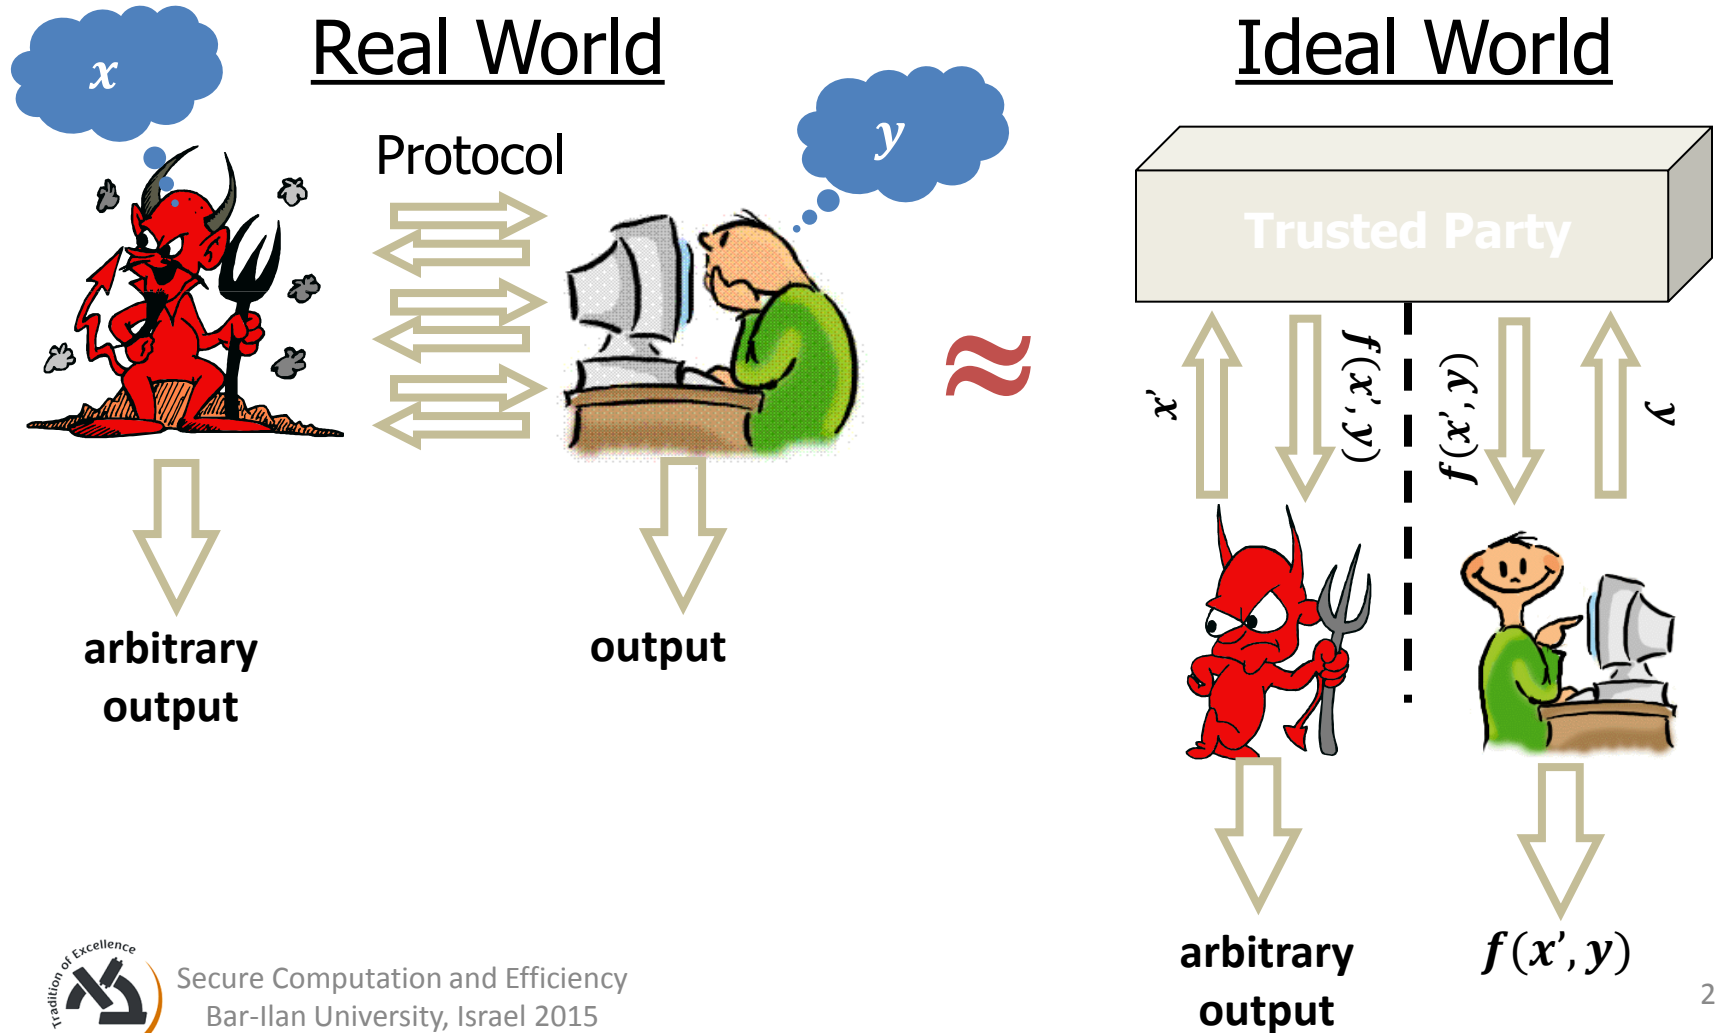
- **What is the best we could hope for?**
 - An incorruptible trusted party
 - All parties send inputs to trusted party (over perfectly secure communication lines)
 - Trusted party computes output
 - Trusted party sends each party its output (over perfectly secure communication lines)
 - This is an **ideal world**
- **What can an adversary do?**
 - Just choose its input...



The Ideal/Real Paradigm

- The real protocol must be like the ideal world
- Formalizing this notion:
 - For *every adversary A* attacking the real protocol, *there exists an adversary S* in the ideal model such that the output distributions (of all) are computationally indistinguishable
 - S simulates a real protocol execution while interacting in the ideal world
 - Here we always look at the joint output distribution

The Ideal/Real Paradigm



“Formal” Security Definition

- **Protocol π securely computes a function f if:**
 - For every non-uniform polynomial-time real-model adversary A , there exists a non-uniform polynomial-time ideal-model adversary S , such that for all input vectors and auxiliary inputs:
 - the joint outputs of A and the honest parties in a real execution of π is computationally indistinguishable from the joint outputs of S and the honest parties in an ideal execution where the trusted party computes f

Properties

- **The following properties hold**
 - **Privacy**: from adversary's outputs
 - **Correctness**: from honest parties' outputs
 - **Independence of inputs**: from ideal execution
 - **Fairness and guaranteed output delivery**: from ideal execution
- More?

Relaxing the Ideal Model

- In some cases, this ideal model is too strong and cannot be achieved
- Fairness cannot be achieved in general without an honest majority

Relaxing the Ideal Model

- **Change the instructions of the trusted party**
 - Trusted party receives input from all parties
 - Trusted party sends corrupted parties' outputs to adversary
 - Adversary says “continue” or “halt”
 - If “continue”, trusted party sends output to honest parties; else, it sends “abort”

Reactive Functionalities

- **Functionalities that obtain inputs and provide outputs in stages**
- **Examples:**
 - Mental poker
 - Commitment schemes
- **This is also useful for relaxing ideal functionalities (give side information to S)**
- **The definition extends naturally to this as well**

Advantages of This Approach

- **General – it captures ALL applications**
- **The specifics of an application are defined by its functionality, security is defined as above**
- **The security guarantees achieved are easily understood (because the ideal model is easily understood)**
 - We can be confident that we did not “miss” any security requirements

Using Secure Computation

- **The ideal-model paradigm**
 - You don't need to understand anything about how a protocol works to use it
 - You just need to imagine an incorruptible trusted party computing the functionality for you
- **Very advantageous for usage**

Sequential Modular Composition

- **Sequential modular composition:**
 - Secure protocols are run sequentially, with arbitrary messages sent in between them
- **Why consider this?**
 - An important security goal within itself
 - Very helpful (if not crucial) tool for analyzing the security of protocols
- **Formalization – Hybrid Model**
 - A trusted party helps to compute a sub-functionality
 - REAL messages & IDEAL messages

Sequential Modular Composition

- Subprotocols ρ_i securely compute functionalities f_i
- Protocol π securely computes g in a hybrid model where a trusted party is used to compute every f_i
 - This is much easier to analyze since each f_i is effectively “perfectly secure”
- **Theorem:** assuming the above, the real protocol π^ρ that uses real calls to each ρ_i instead of a trusted party for f_i , securely computes g .

Concurrent Composition

- **We have considered the stand-alone model**
 - This implies sequential composition
- **What about concurrent composition?**
 - An Internet-like setting where many (arbitrary, secure and insecure) protocols are run concurrently, with the adversary controlling the scheduling
- **This models the real-world setting more accurately**
 - We don't know what the result is of running stand-alone protocols concurrently with related inputs

Concurrent Composition

- **Concurrent general composition**
 - Strictly harder than the stand-alone model
 - *Impossible* without some trusted set-up assumption (like a common reference string)
- **The UC definition (universal composability) guarantees security in this setting**
 - Efficient UC security is a special challenge...
- **Recommended to study UC next, after studying the stand-alone setting**

Relaxed Definitions

- **In order to achieve high efficiency, sometimes can consider weaker definitions**
 - Semi-honest (but this is very weak)
 - Covert adversaries: adversary may be malicious but is guaranteed to be caught cheating with good probability
 - Suitable where adversaries can be penalized for being caught cheating (e.g., business loss)
 - Privacy only (malicious)
 - Problematic...

Summary

- **Semi-honest: simulator given input/output generates the adversary's view**
 - Probabilistic functionalities – must consider joint distribution of view and outputs
 - Deterministic functionalities: easier, suffices to separately consider correctness and view simulation
- **Malicious: ideal-real simulation**
- **Sequential composition**
- **Advanced topics**
 - Concurrent composition
 - Relaxed definition
 - Semi-honest vs malicious

General vs Specific Protocols

- **Most of the school will focus on general protocols**
 - Convert the function into a Boolean or arithmetic circuit
 - Compute the circuit securely
- **It seems that for specific problems, specific protocols should be more secure**


General vs Specific Protocols

- **General protocols – advantages**
 - Implement once
 - Very flexible: almost no difference between
 - Set intersection
 - Size of set intersection
 - Output 1 if set intersection size is greater than k
 - In many cases is competitive, and in fact the fastest solution known

OBLIVIOUS TRANSFER



Oblivious Transfer (OT)



Called 1-out-of-2 oblivious transfer (OT_1^2)

Fundamental Primitive

- OT is **complete**
 - If can compute OT then can compute any functionality
- **Constructing OT**
 - OT cannot be constructed from PKE in a black box manner
 - Can be constructed from
 - Enhanced trapdoor permutations
 - DDH, RSA, Lattices

Just a Few Important OT Results

- OT is symmetric
- Can construct efficient OT_1^N and OT_k^N from OT_1^2
- Can construct malicious OT from semi-honest OT in a black-box manner (inefficiently)
- Many variants of OT are equivalent
 - Random OT
 - Rabin OT
 - Weak OT

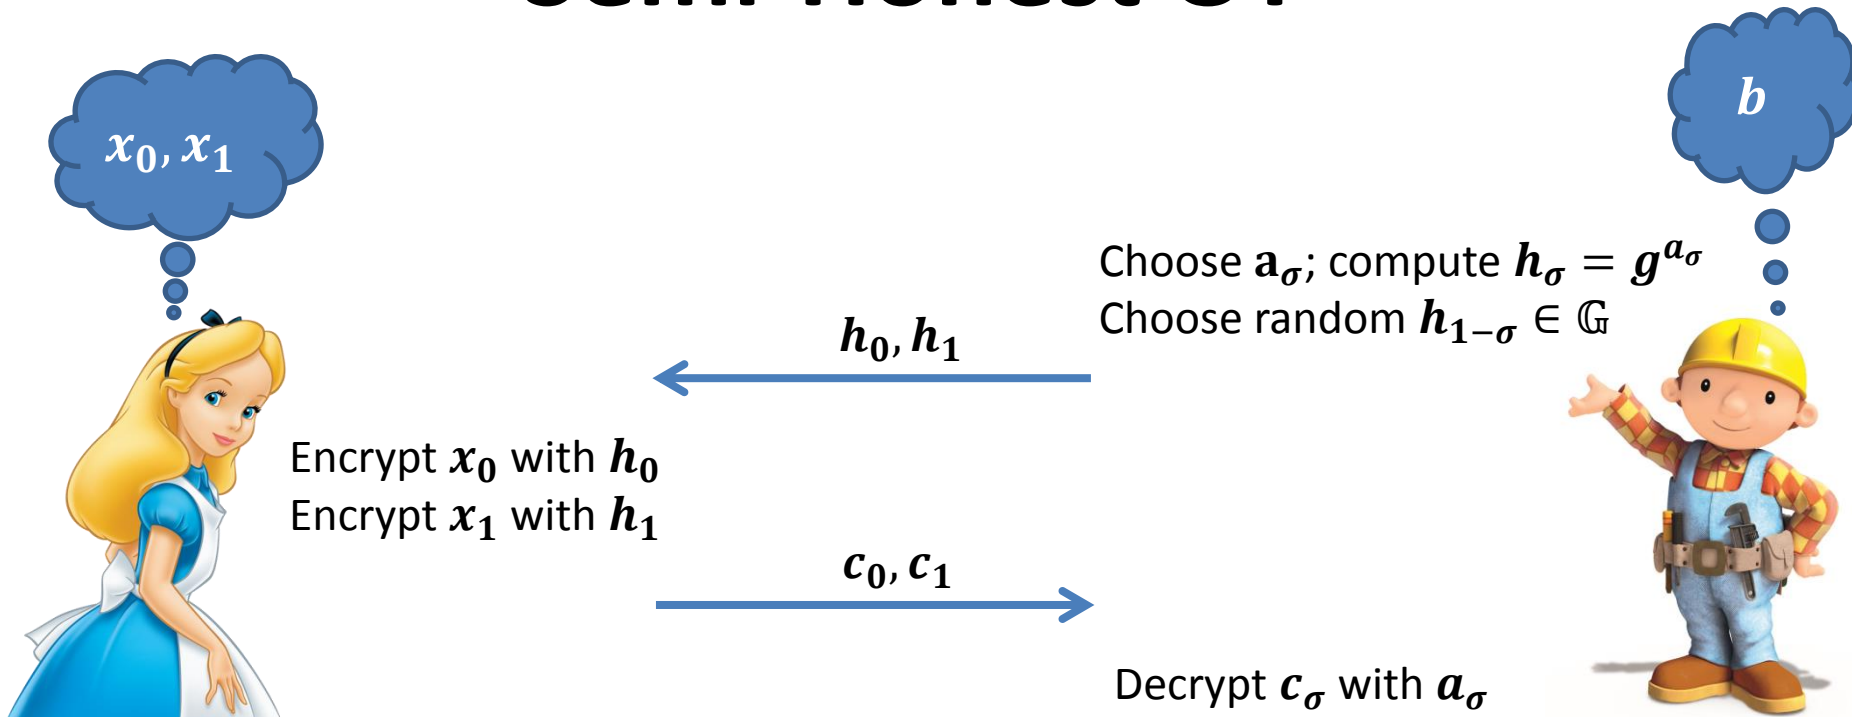
Efficient OT from DDH

- **Recall the DDH assumption over a group \mathbb{G} of order q with generator g**
 - The DDH assumption says that
$$\{(g, g^a, g^b, g^{ab})\} \approx \{(g, g^a, g^b, g^c)\}$$
where $a, b, c \leftarrow \mathbb{Z}_q$ are random

Semi-Honest OT

- **Recall ElGamal encryption**
 - **Secret key:** random $a \leftarrow \mathbb{Z}_q$
 - **Public key:** $h = g^a$
 - **Encrypt** $m \in \mathbb{G}$: $c = (u, v) = (g^r, h^r \cdot m)$, random $r \in \mathbb{Z}_q$
 - **Decrypt** (u, v) : compute $m = \frac{v}{u^a}$
 - Note: $\frac{v}{u^a} = \frac{h^r \cdot m}{(g^r)^a} = \frac{h^r \cdot m}{(g^a)^r} = \frac{h^r \cdot m}{h^r} = m$

Semi-Honest OT



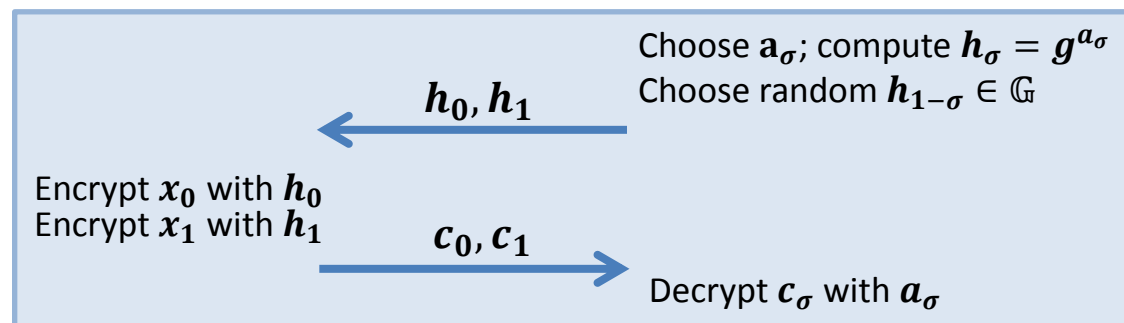
Note:

- Encrypt x_0 with h_0 : $(u_0, v_0) = (g^r, (h_0)^r \cdot x_0)$
- Encrypt x_1 with h_1 : $(u_1, v_1) = (g^s, (h_1)^s \cdot x_1)$

Semi-Honest OT – Security

- **Security:**

- Alice sees only two public keys, which are two random group elements (and so learns nothing about σ)
 - Formally, simulate by sending two random group elements
- Bob knows only one private key and so learns only x_σ
 - Formally, simulate by encrypting x_σ with h_σ , and encrypting garbage (e.g., 0) with $h_{1-\sigma}$



More Efficient Semi-Honest OT




Choose $r \leftarrow \mathbb{Z}_q$
Compute $u = g^r$
Compute $v_0 = (h_0)^r \cdot x_0$
Compute $v_1 = (h_1)^r \cdot x_1$

h_0, h_1

u, v_0, v_1

Choose a_σ ; compute $h_\sigma = g^{a_\sigma}$
Choose random $h_{1-b} \in \mathbb{G}$



Output $x_\sigma = \frac{v_\sigma}{u^{a_\sigma}}$

Malicious Adversaries

- **Corrupted sender:**
 - Sender cannot cheat
 - Simulator can “extract” both x_0, x_1 by choosing both h_0 and h_1 so that it knows the secret keys
- **Corrupted receiver:**
 - Receiver can choose both h_0 and h_1 so that it knows the secret keys

Preventing Malicious

- **The idea:**
 - Alice sends a random group element w
 - Bob chooses h_0, h_1 so that $h_0 \cdot h_1 = w$
 - Bob can easily do this by choosing a_σ , computing $h_\sigma = g^{a_\sigma}$ and setting $h_{1-\sigma} = w/h_\sigma$
 - Bob cannot know both DLOGs of h_0, h_1 or it can compute the DLOG of H
- **Encryption uses a random oracle since “not completely knowing” a secret key doesn’t suffice**
 - Encrypt by $(g^r, \text{HASH}((h_0)^r) \oplus x_0), \dots$

State of the Art – OT

- **Semi-honest adversaries**
 - Receiver: 2 exponentiations + send 2 group elements
 - Sender: 3 exponentiations + send 3 group elements
- **Malicious adversaries (Random Oracle)**
 - Same as semi-honest
- **Malicious adversaries (PVW)**
 - Receiver: 3 exponentiations + send 2 group elements
 - Sender: 8 exponentiations (effectively 6) + send 4 group elements

Proving Malicious Security

- Proving security in the malicious model is tricky and subtle
- The ideal/real model paradigm
 - Need a simulator who internally runs the real adversary and externally interacts with the trusted party (sending input and getting output)
 - The simulator needs to “**extract**” the real adversary’s input, get output, and make the output match
- We demonstrate the ideal/real proof technique for the problem of **coin tossing**

Proving Malicious Security

- **Blum's protocol (with ElGamal):**
 - Party P_1 :
 - Choose random $b \in \{0,1\}$ and $r, s \leftarrow \mathbb{Z}_q$
 - Compute $h = g^r$, $u = g^s$, $v = h^s \cdot g^b$
 - Send (h, u, v) to P_2
 - Party P_2 :
 - Choose random $b' \in \{0,1\}$
 - Send b' to P_1
 - Party P_1 sends r, s, b to P_2
 - Party P_2 verifies that $h = g^r$, $u = g^s$, $v = h^s \cdot g^b$
 - Both parties output $b \oplus b'$

Intuition

- **Consider a corrupt P_2**
 - By the security of El Gamal encryption, it knows nothing about b when it chooses b'
- **Consider a corrupt P_1**
 - The values (h, u, v) fully define b
 - There exists a single pair (r, s) so that $h = g^r, u = g^s$
 - The value v can either be h^s or $h^s \cdot g$, but **not both**
 - P_2 chooses b' after P_1 sends b ; by the above, P_1 cannot change b and so P_1 cannot bias the output

Proving Security - P_1 corrupted

- Let A be an adversary; S works as follows
- S receives a random bit β from the trusted party
- S invokes A and receives (h, u, v)
- S works as follows:
 - S *internally* hands A the value $b' = 0$
 - S **rewinds** A and internally hands A the value $b' = 1$
 - If A replies correctly both times, S learns the value b , sets $b' = b \oplus \beta$, and outputs this as A's view. In addition, A *externally* sends **continue** to the TTP
 - If A does **not** reply correctly either time, S sends **abort** to the TTP and outputs a random b' as A's view
 - If A aborts once, then S learns the value b , sets $b' = b \oplus \beta$, and outputs this as A's view. If A aborts on this b' then S sends **abort** to the TTP; else it sends **continue** to the TTP

Proving Security - P_2 corrupted

- Let A be an adversary; S works as follows
- S receives a random bit β from the trusted party
- S invokes A and works as follows:
 - S chooses a random b and internally hands A the tuple (h, u, v) computed correctly for b
 - S receives b' from A
 - If $b \oplus b' = \beta$ then S outputs (h, u, v) and (r, s, b) as its view, and sends **continue** to the TTP
 - Else, S **rewinds** A and goes to the beginning again
- Note: there is no abort here since we can just take $b' = 0$ as default if P_2 doesn't respond

Summary

- **Oblivious transfer is a fundamental primitive**
 - It is heavily used in most general secure computation protocols
- **Oblivious transfer is very efficient**
 - But it does cost exponentiations every time!
 - This afternoon we will see how to improve this